

Evaluating Differential Privacy in Federated Clustering Under Adversarial Attacks

Presented by:

Urvi Gupta (22B1006) Dept. of CSE

Ojas Maheshwari (22B0965) Dept. of CSE

Course: CS6007:Multi-Agent Machine Learning

Instructor: Prof. Avishek Ghosh

IIT Bombay

November 2, 2025

Problem Statement

- **Core challenge:**
 - Evaluate whether Differential Privacy can effectively protect Federated K-Means from adversarial attacks **without significantly degrading clustering quality.**
- **Key difficulties:**
 - Adding DP noise protects privacy but degrades clustering quality — finding a meaningful trade-off is challenging.
 - Simulating realistic and effective attacks that meaningfully test privacy guarantees is non-trivial.
 - Non-IID client data and varying sample sizes complicate both clustering and privacy analysis.
- **What prior work misses:**
 - Prior work focuses on theoretical privacy guarantees (ϵ, δ) but rarely tests robustness against concrete attacks.
 - Empirical evaluation under realistic attack settings is largely missing, leaving practical effectiveness uncertain.

Literature Survey - Federated Clustering

- **Federated K-Means Clustering (Garst et al., Delft University):**
 - Proposed one of the first frameworks for performing K-Means clustering in a federated setting without sharing raw data.
 - Highlighted communication efficiency and privacy concerns, but did not incorporate formal privacy mechanisms.
- **Differentially Private Federated K-Means (Scott et al., 2025):**
 - Introduced a differentially private version of federated K-Means by adding calibrated Gaussian noise on server-side updates.
 - Provided theoretical privacy guarantees (ϵ, δ) and utility analysis, but no empirical evaluation against real attacks.
- **Inverting Gradients (Geiping et al., 2020):**
 - Showed that gradients shared during federated learning can be reversed to recover the original training data.
 - Highlighted serious privacy risks in model updates, motivating similar reconstruction studies in federated clustering.

Contribution

- **Verification and Benchmarking:** Verified the FedDP-KMeans implementation of **Scott et al. (2025)** against the non-private Federated K-Means baseline.
- **Reconstruction attacks:**
 - **Data-level reconstruction:** Simulated recovery of an individual datapoint from shared client updates, comparing reconstruction quality (cosine similarity) under DP vs non-DP.
 - **Client-level reconstruction:** Attempted to reconstruct the client's aggregate update (cluster means sent to the server) to evaluate information leakage at the client level.

Definitions

Evaluating K-Means Performance

How we measure success

1. K-Means Cost

Definition: It is the **sum of squared Euclidean distances** from each data point to its assigned cluster center. Given a dataset

$P = \{p_1, \dots, p_N\}$, a set of K centers $C = \{c_1, \dots, c_K\}$, and assignments y_i for each point p_i : $\text{Cost}(P, C) = \sum_{i=1}^N \|p_i - c_{y_i}\|^2$

2. K-Means Accuracy

- ① For each predicted cluster j , find the *majority* true label among all points assigned to j .
- ② Count how many points in cluster j have this majority true label. This is the number correct for cluster j .
- ③ Sum these counts across all K clusters to get the total number correct.
- ④ Accuracy =
$$\frac{\text{Total Number Correct}}{\text{Total Data Points}}$$

Differential Privacy (DP)

Definition

An algorithm $A : \mathcal{P} \rightarrow \mathcal{S}$ is said to satisfy (ε, δ) -**Differential Privacy** if for all neighboring datasets P, P' and all measurable subsets $S \subseteq \mathcal{S}$,

$$\Pr[A(P) \in S] \leq e^\varepsilon \Pr[A(P') \in S] + \delta$$

where $\varepsilon > 0$ quantifies the privacy loss and $\delta \geq 0$ allows for a small probability of failure.

- **Neighboring Datasets:**
 - *Data-point-level DP*: P and P' differ by one individual data point.
 - *Client-level DP*: P and P' differ by all data points belonging to one client.
- **Compositionality**: If algorithms A_1, \dots, A_t are each $(\varepsilon_i, \delta_i)$ -DP, their combination remains DP with parameters $(\sum_i \varepsilon_i, \sum_i \delta_i)$, enabling multi-step private pipelines.

Privacy Models (1/2): Data-point Privacy

- **What it Protects:**
 - A *single data point* (e.g., one person's record).
- **The Privacy Question:**
 - "Does the final result change if we add or remove *Jane Doe's record* from Client A's dataset?"
- **Sensitivity (L_2):**
 - The maximum impact of **one point**.
 - This is a small and generally known value (or easy to set).
 - Example: `max_data_norm`
- **Clipping:**
 - Applied to each data point (or implicitly assumed).
- **Use Case:**
 - Standard "cross-silo" FL.
 - Clients are trusted entities (like hospitals).
 - The individuals *within* the silo need privacy.

Privacy Models (2/2): Client-level Privacy

- **What it Protects:**
 - The *entire contribution* of one client (e.g., all 10,000 records from "Hospital A").
- **The Privacy Question:**
 - "Does the final result change if we add or remove *Hospital A's entire dataset* from the training?"
- **Sensitivity (L_2):**
 - The maximum impact of **one client**. This is large and unknown.
 - Clipping is mandatory to create a finite bound.
 - Example: `max_client_norm_vector`
- **Clipping:**
 - Applied to the client's *total aggregated update* before it leaves the client.
- **Use Case:**
 - Standard "cross-device" FL.
 - The client (e.g., a mobile phone) is the user and cannot be trusted.
 - This is a much stronger privacy guarantee.

Federated Differentially Private K-Means (FedDP-KMeans)

- **Goal:** Perform k -means clustering across m clients without sharing raw data, while ensuring (ε, δ) -differential privacy.
- **Setup:** Each client j has data P_j , and the server has small auxiliary data Q (may be out-of-distribution).
- **Algorithm Structure:**
 - ① **FedDP-Init:** Generates DP-protected initial cluster centers using client updates and server data.
 - ② **FedDP-Lloyds:** Refines centers iteratively with noised client contributions.
- **Privacy:** Gaussian noise added to client-side aggregates.

FedDP-Init — Algorithm

- **Objective:** Generate high-quality, privacy-preserving initial cluster centers before running Lloyd's iterations.

- **Step 1 — Subspace Estimation:**

- Each client computes its local data covariance $P_j P_j^T$.
- The server aggregates these (with Gaussian noise) to estimate top- k eigenvectors.
- Projects data into a lower-dimensional subspace, improving signal-to-noise ratio.

- **Step 2 — Proxy Weight Computation:**

- Server shares projection matrix Π and projected server data ΠQ .
- Clients compute how many local points are closest to each server point (weights).
- Noised weights are sent to server (Laplace mechanism, $(\varepsilon_2, 0)$ -DP).

- **Step 3 — Center Refinement:**

- Server runs weighted k -means on ΠQ to get projected centers.
- Clients refine these using their local data and send noised sums/counts.
- Server aggregates and reconstructs final DP initialization centers.

FedDP-Lloyds Algorithm

- **Goal:** Refine cluster centers under DP after FedDP-Init.
- **Key Idea:** Lloyd's updates rely only on *sums and counts* — enabling secure aggregation and DP noise addition.
- **Steps per Round:**
 - Server sends current centers ν^{t-1} to clients.
 - Client j assigns each point to the nearest center and computes:

$$m_{jr} = \sum_{p \in S_r^j} p, \quad n_{jr} = |S_r^j|$$

- Add Gaussian noise to m_{jr} and Laplace noise to n_{jr} .
- Server aggregates noisy stats and updates:

$$\nu_r^t = \tilde{m}_r / \tilde{n}_r$$

- **Privacy:** Overall algorithm achieves (ε, δ) -DP via composition.

Explanations and Intuition

Lloyd's Algorithm & The Importance of Initialization

- **The Problem with Lloyd's in DP:**

- The paper identifies a counterintuitive trade-off for DP algorithms.
- Accessing private data more often (e.g., more rounds of Lloyd's algorithm) requires more noise per step to stay within the privacy budget.
- This can lead to *lower* accuracy, not higher.

- **The Goal:**

- Consequently, a good initialization is crucial for achieving high accuracy.
- This is because it allows the algorithm to avoid running many iterations of Lloyd's.

- **Iterations Required for FedDP-KMeans:**

- **Finding:** The FedDP-Init centers are so close to the optimal ones that only very few (sometimes none at all) steps of Lloyd's algorithm are required to refine them.
- **Conclusion:** The preference was to instead use all the budget for the initialization rather than on subsequent Lloyd's steps.

Why FedDP-Init Requires Fewer Lloyd's Steps?

More iterations → smaller budget per step → **more noise per step**.

- **k-FED: A "Local-First" Initialization**

- k-FED (a non-private baseline) has clients run *k*-means *locally* and send their centers.
- If client data is non-IID (unbalanced or different), these local centers are **highly biased**.
- The server then clusters these biased centers, resulting in a poor global starting point that is far from the true optimum.

- **FedDP-Init: A "Global-First" Initialization**

- FedDP-Init is designed to solve this. It uses its budget to get a high-quality start.
- **Step 1 (Subspace)** estimates the *global* data structure from all clients.
- **Step 3 (Refinement)** is already one complete, *globally-informed* Lloyd's step.
- **Result:** The centers are already so close to the optimum that few (or even **zero**) further noisy FedDP-Lloyds iterations are needed.

Clipping: An Important Step in FedDP-Init

What is Clipping? (What, Why, How)

- **What is it?**

- Clipping enforces a **maximum limit** (a "cap") on the size of a client's (or datapoint's) contribution before it is sent to the server.

- **Why is it necessary?**

- Differential Privacy adds noise based on **sensitivity** (the max possible impact of one client).
- A client with 1M points has a *massive* impact vs. a client with 10 points. This sensitivity is unknown and potentially infinite.
- Clipping creates a **fixed, known sensitivity** (e.g., "no client can contribute more than 100.0"). This makes adding a finite amount of noise possible.

- **How is it implemented? (L_2 Clipping)**

- The client calculates the L_2 norm (length) of its update vector (e.g., its local sum m_r).
- **If $\text{norm} \leq C$:** The vector is sent as-is.
- **If $\text{norm} > C$:** The vector is "shrunk" by multiplying it by C / norm . Its original direction is thus preserved.

Parameters of FedDP-Init: Clipping Bounds

- `outer_product_clipping_bound`: Caps the client's **covariance matrix** ($P_j P_j^T$) for Step 1 of FedDP-Init.
- `weighting_clipping_bound`: Caps the client's **proxy weights** (counts) for Step 2 of FedDP-Init.
- `center_init_clipping_bound`: Caps the L_2 norm of the client's **local mean vectors** for Step 3 of FedDP-Init.
- `center_init_contributed_components_clipping_bound`: Caps the L_1 norm (count) of the client's **histogram of non-empty clusters** for Step 3.
- `fedlloyds_clipping_bound`: Caps the L_2 norm of the client's **vector sums** (m_r) during FedLloyds rounds.
- `fedlloyds_laplace_clipping_bound`: Caps the L_1 norm (count) of the client's **point counts** (n_r) during FedLloyds rounds.

Comparative Analysis of KFed and FedDP-Init

Datasets Used

- **Synthetic Gaussian Dataset:**

- A controlled, high-dimensional dataset used to simulate a standard federated k -means environment.
- Ideal for isolating the effect of Differential Privacy (DP) on clustering utility and privacy leakage due to its simple, predictable structure.


- **Folktale Dataset (US Census ACS):**

- A real-world dataset focused on demographic and economic features.
- Typically lower-dimensional and **non-isotropic**, providing a practical test against complex, real-world data distributions.

Gaussian Synthetic Dataset with Datapoint Privacy


Comparison of Utility vs. Privacy between KFed and FedDP-Init

Comparative Privacy-Utility Tradeoff (Gaussian Data-Point)



Gaussian Synthetic Dataset with Datapoint Privacy

Privacy-Utility Tradeoff Analysis (Gaussian Data-Point)




Why FedDP-Init Won with Data-Point Privacy

- **Scenario:** Low noise (data-point level).
- **FedDP-Init**
 - Runs 3 private steps to get its initialization.
 - Because the data-point noise is tiny, even 3 steps of this noise have almost no effect.
 - It produces a near-perfect initialization from the start.
- **KFed:**
 - Gets a free (non-private) initialization, which is not as good as the one FedDP-Init finds.
 - It never catches up to FedDP-Init's superior starting point.
- **Conclusion:** With low noise, the **quality of the initialization algorithm** is what matters most. The paper's FedDP-Init is simply a better algorithm.

Gaussian Synthetic Dataset with Client-Level Privacy

Comparison of Utility vs. Privacy between KFed and FedDP-Init

Comparative Privacy-Utility Tradeoff (Gaussian Client-Level)




Why KFed Wins at Low Epsilon (Client-Level Privacy)

- **Scenario:** High noise (client-level).
- **FedDP-Init:**
 - At low epsilon (e.g., 0.25), it splits its tiny budget 3 ways for its 3 private steps.
 - Each step must add massive client-level noise.
 - This triply large dose of noise destroys the initialization, resulting in terrible accuracy and cost.
- **KFed:**
 - Gets its free non-private initialization.
 - It adds the massive client-level noise *only once* during refinement.
 - One dose of massive noise is less bad than three, so it performs better at low epsilon.
- However, at $\epsilon = 2.5$, FedDP-Init gave an accuracy of 97.62% that is much higher than what KFed could ever achieve. Thus, the new algorithm eventually surpasses KFed at higher epsilon.
- **Conclusion:** With high noise, the **number of private steps** is what matters most. FedDP-Init's 3-step process is a disadvantage at low epsilon.

Folktables US Dataset with Datapoint Privacy

K-Means Cost vs. Epsilon



Reconstruction Attacks

Threat Model

- **White-box setting:** full access to training code and DP mechanism.
- Observes global cluster centers after each round.
- knows the DP mechanism (Gaussian/Laplace) and its parameters (ε, δ , noise scales, clipping bounds).
- Attacker knows number of clients, number of features, and the data distribution/statistics (but **not** raw datapoints).
- With knowledge of model hyperparameters and training setup, attacker can train **shadow models** to judge whether a point (or client) was likely included.

Note: This is a strong (worst-case) attacker model used to stress-test DP; a weaker black-box attacker that only sees final centers would yield much weaker attacks and could give a false sense of robustness.

Reconstruction Attack: Overview

- **Goal:** Reconstruct **private training information** from shared, noisy client updates in Federated Learning (FL).
- **Setup:** The attacker observes the noisy client updates (total sums and counts in Federated K-Means) in each FL round.
- **Attack Levels:**
 - ▷ **Data-Level:** Recover an individual datapoint x .
 - ▷ **Client-Level:** Reconstruct a client's mean update μ_j .
- **Comparison:** Attack success is measured on **Non-Private** vs. **DP-Federated K-Means** algorithms via L2 error and cosine similarity.

Data-Level Reconstruction: Reconstructing \mathbf{x}

- **Objective:** Recover a single, targeted datapoint \mathbf{x} .
- **How the attack works:**
 - ① Run federated training *once without* \mathbf{x} to get baseline global centers (C_{base}).
 - ② Using C_{base} , simulate two noisy updates from the target client:
 - \tilde{u}_{in} : update *with* the target point \mathbf{x} .
 - \tilde{u}_{out} : update *without* the target point \mathbf{x} .
 - ③ Compute the **difference** (the point's isolated contribution):

$$\Delta = \tilde{u}_{\text{in}} - \tilde{u}_{\text{out}}$$

- **Reconstruction Formula:**
 - Find the point's most probable cluster k^* (where the count difference Δ_{counts} is largest): $k^* = \operatorname{argmax}_k (\Delta_{\text{counts}}[k])$
 - The reconstructed point is the **sum difference from that single cluster**:

$$\hat{\mathbf{x}} = \Delta_{\text{sums}}[k^*]$$

- **Metric:** Accuracy is measured by **Cosine Similarity** $\left(\frac{\mathbf{x} \cdot \hat{\mathbf{x}}}{\|\mathbf{x}\| \|\hat{\mathbf{x}}\|} \right)$.

Client-Level Reconstruction: Reconstructing $\hat{\mu}_j$

- **Objective:** Reconstruct the **mean update** ($\hat{\mu}_j$) of a target client j 's entire dataset.
- **How the attack works:**
 - ① Simulate a global model state trained *without* the target client j .
 - ② Simulate the target client j 's **full noisy local update** (\tilde{u}_j).
 - ③ **The noisy update IS the reconstructed total** (sum and count) contribution:

$$(\hat{\Sigma}_j, \hat{N}_j) = \tilde{u}_j \quad (\text{Noisy Sum and Count})$$

- **Reconstructed Mean:** Calculate the mean from the noisy components:

$$\hat{\mu}_j = \frac{\hat{\Sigma}_j}{\hat{N}_j}$$


- **Privacy Effect:** Stronger Differential Privacy (smaller ε) \rightarrow higher **noise distortion** in $\hat{\mu}_j$ \rightarrow lower attack accuracy.

Experiment 1

Data-point Reconstruction Attack with Gaussian Dataset

- **Goal:** Reconstruct the vector of a single target data point from a client's noisy update.
- **Model:** Federated k-means with and without differential privacy
- Clients (M): 100
- Samples per client: 1000
- Feature Size (D): 100
- Clusters (K): 10
- Attack Trials: 50 target data points
- Privacy level: data-point
- Target Epsilon (ϵ): 1.0
- Target Delta (δ): 1.0×10^{-6}
- Core Clipping (C): 11 (FedLloyds L₂ Clipping Bound)
- Mechanism: Gaussian Noise (due to L₂ clipping and $\delta \neq 0$)

Experiment 1 - Results




Experiment 2

Client Level Reconstruction Attack with Gaussian Dataset

- **Goal:** Reconstruct a client's mean update
- **Model:** Federated k-means with and without differential privacy
- Clients (M): 100
- Samples per client: 1000
- Feature Size (D): 100
- Clusters (K): 10
- Attack Trials: 20 clients
- Privacy level: client-level
- Target Epsilon (ϵ): 1.0
- Target Delta (δ): 1.0×10^{-6}
- Core Clipping (C): 120 (FedLloyds L₂ Clipping Bound)
- Mechanism: Gaussian Noise (due to L₂ clipping and $\delta \neq 0$)

Experiment 2 - Results




Experiment 3

Data-point Reconstruction Attack with Folktale Dataset

- **Goal:** Reconstruct the vector of a single target data point from a client's noisy update.
- **Model:** Federated k-means with and without differential privacy
- Clients (M): 51
- Samples per client: Variable
- Feature Size (D): 124
- Clusters (K): 10
- Attack Trials: 10 target data points
- Privacy level: data-point
- Target Epsilon (ϵ): 1.0
- Target Delta (δ): 1.0×10^{-6}
- Core Clipping (C): 2.65 (FedLloyds L₂ Clipping Bound)
- Mechanism: Gaussian Noise (due to L₂ clipping and $\delta \neq 0$)

Experiment 3 - Results



Conclusion: Key Takeaways & Challenges

- **Key Takeaways:**

- Spending privacy budget on initialization (FedDP-Init) is more effective than on iterations (KFed/FedLloyds).
- FedDP-Init provides high accuracy with very few (or even 0) Lloyd's steps.
- FedDP-Init is highly robust to reconstruction attacks with minimal utility loss.

- **Limitations & Challenges:**

- Client-level Folktables simulation failed due to clipping issues and sparse data.
- Membership Inference Attacks (MIA) were inconclusive, succeeding similarly on both private and non-private models.

Future Work

- **Address Current Limitations:**

- Develop systematic methods to estimate optimal hyperparameters (e.g., clipping bounds).
- Investigate and design more effective defenses against Membership Inference Attacks.

- **Explore New Directions:**

- Test the system's robustness against other types of privacy attacks (e.g., attribute inference).
- Extend the DP-KMeans framework to different data modalities or more complex clustering algorithms.

References I

- [1] S. Garst and M. Reinders.
Federated K-Means Clustering.
Delft Bioinformatics Lab, Delft University of Technology, 2023.
- [2] J. Scott, C. H. Lampert, and D. Saulpic.
Differentially Private Federated k-Means Clustering with Server-Side Data.
arXiv preprint arXiv:2506.05408, 2025.
- [3] R. Shokri, M. Stronati, C. Song, and V. Shmatikov.
Membership Inference Attacks Against Machine Learning Models.
In *2017 IEEE Symposium on Security and Privacy (SP)*, pages 3–18, 2017.
- [4] J. Geiping, H. Bauermeister, H. Dröge, and M. Moeller.
Inverting Gradients: How Easy Is It to Break Privacy in Federated Learning?
In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.

Thank You!